

Distributed in Space and Time

Dr. Roland Kuhn
Akka Tech Lead
[@rolandkuhn](https://twitter.com/rolandkuhn)



The Future is Bright



source: <http://nature.desktopnexus.com/wallpaper/58502/>

The Future is Bright

- Computing Resources: Effectively Unbounded
 - 16-core CPU packages, 1024 around the corner
 - on-demand provisioning in the cloud
- Demand for Web-Scale Apps: Ravenous
 - everyone has the tools
 - successful ideas are bought for Giga-\$

How do I make a Large-Scale App?

- figure out a service to provide to everyone
- this will **require resources** accordingly:
 - storage
 - computation
 - communication
- scalability implies **no shared contention points**
- need to be elastic → **use the cloud**

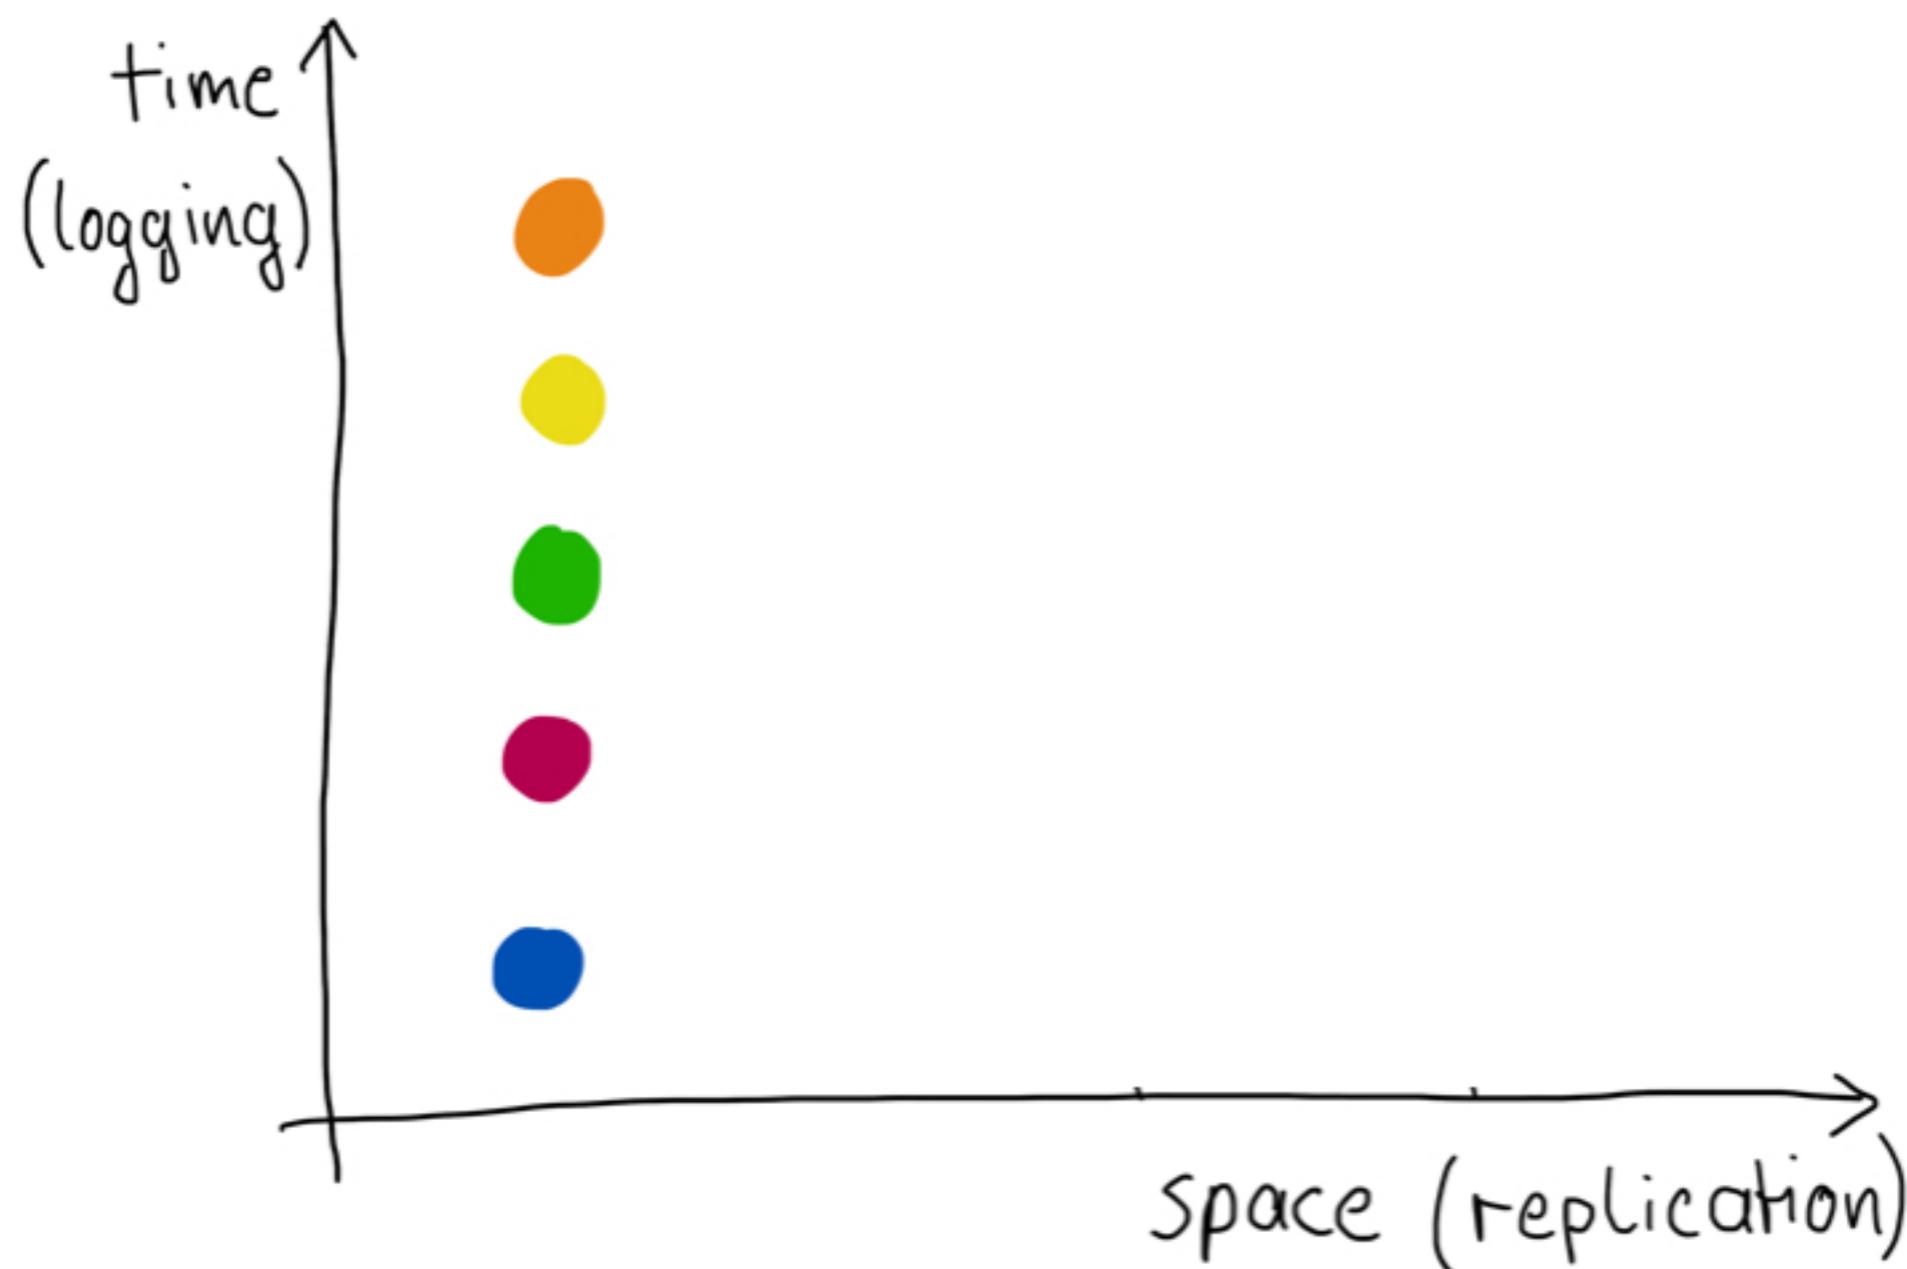
How do I make an Always-Available App?

- must not have single point of failure
- all functions and resources are **replicated**
 - protect against data loss
 - protect against service downtime
 - regional content delivery
- need to distribute → **use the cloud**

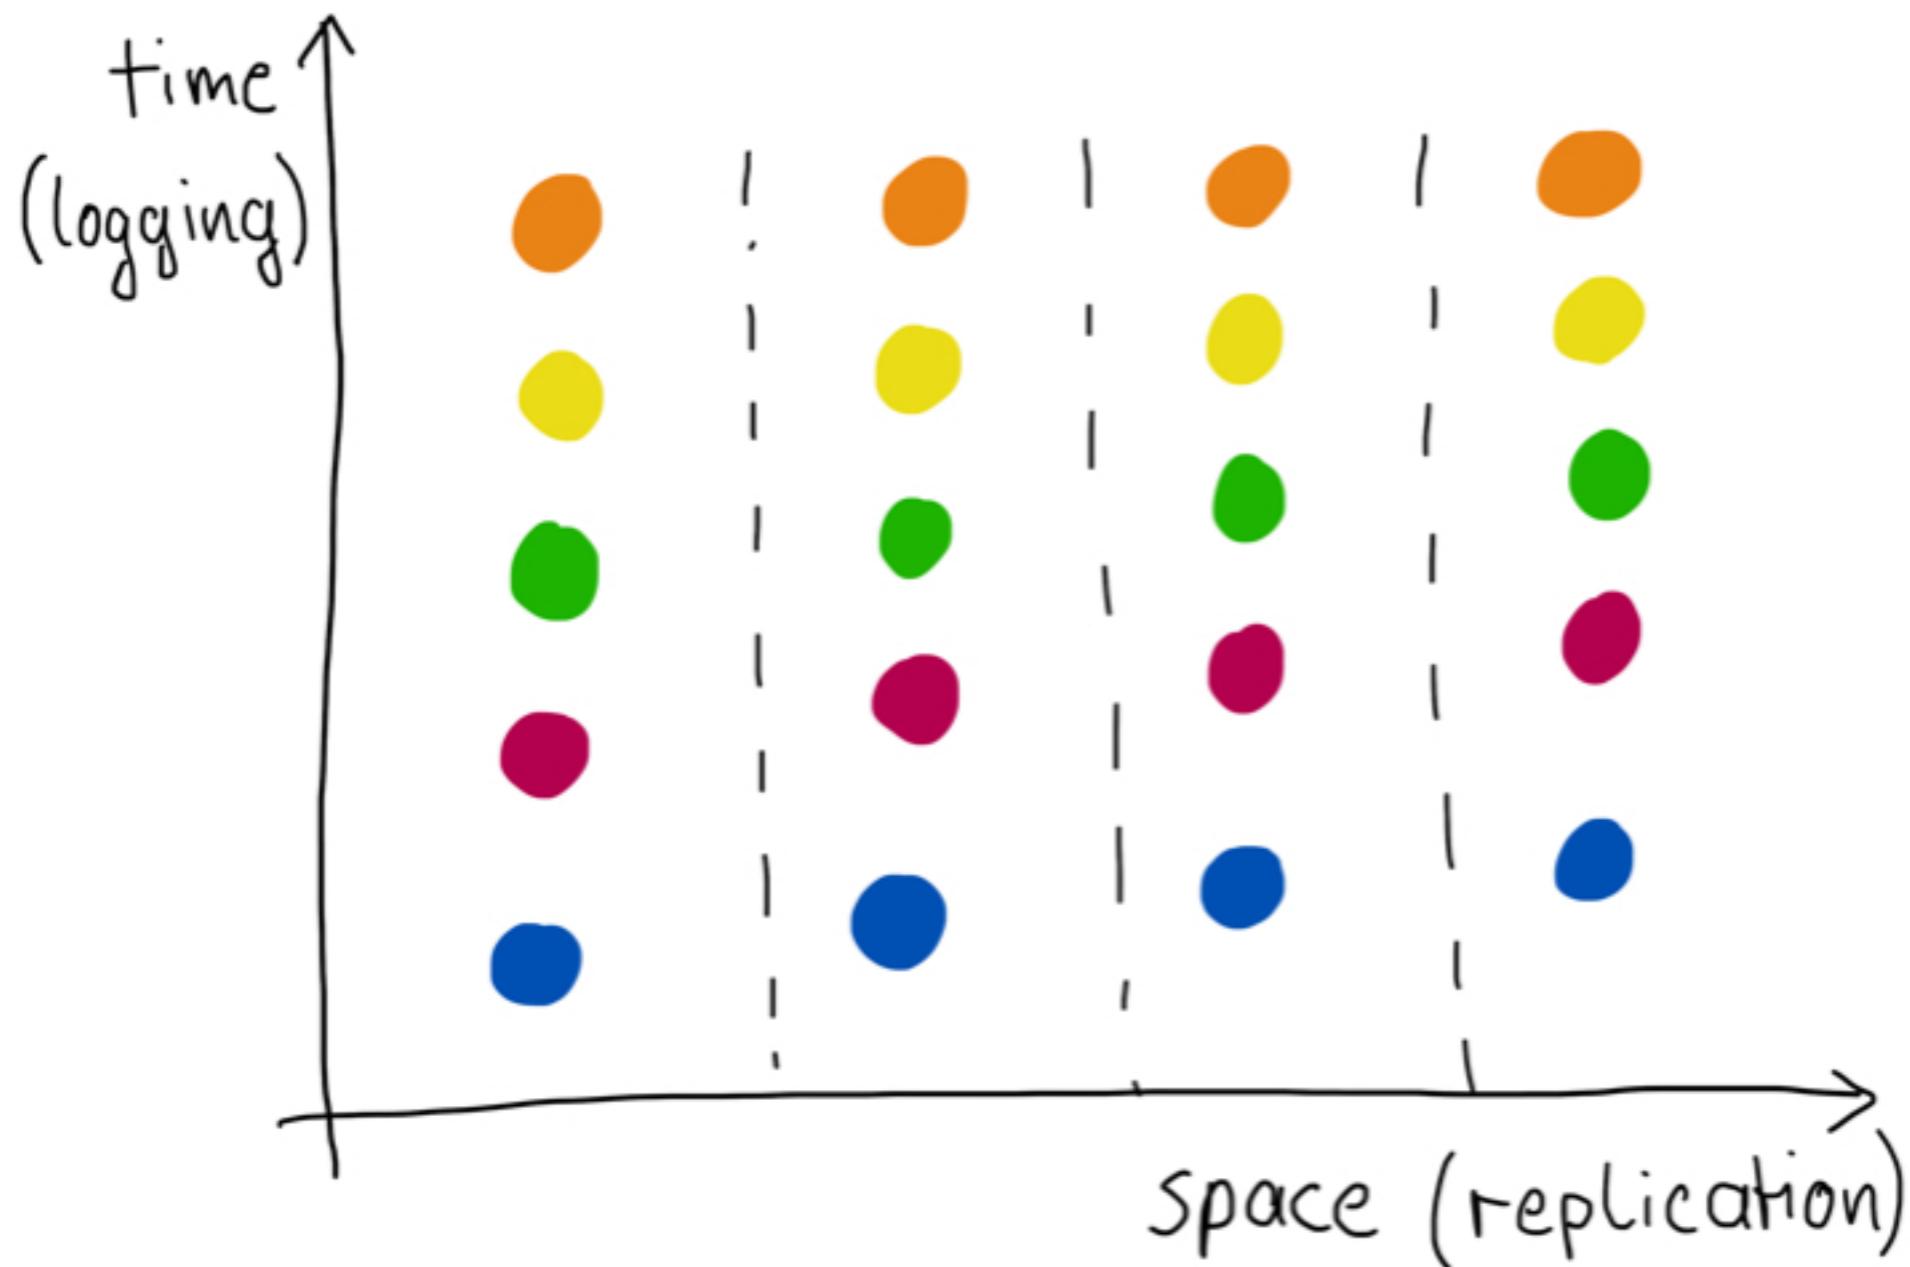
How do I make a Persistent App?

- runtime replication usually not sufficient
 - system of record needed
- one central database is not tenable
- persist “locally” and **replicate changes**

How do I make a Persistent App?



How do I make a Persistent App?



How do I make a Persistent App?

- runtime replication usually not sufficient
 - system of record needed
- one central database is not tenable
- persist “locally” and **replicate changes**
 - temporally distributed change records
 - spatially distributed change log
 - granularity can be one Aggregate Root (Event Sourcing)
 - extract business intelligence from history
- scale & distribute persistence → **in the cloud**

As Noted Earlier: The Future is Bright



source: <http://nature.desktopnexus.com/wallpaper/58502/>

... or is it?



Photograph: Patrick Cromerford, 2011

Distribution: a Double-Edged Sword

- did that message make it over the network?
- is that other machine running?
- are our data centers in sync?
- did we lose messages in that hardware crash?

The Essence of the Problem

A distributed system is one whose parts can fail independently of each other.

Fighting Back!



source: www.hear2heal.com (Kokopelli Rain Dance)

PAXOS and Friends

- problem: distributed consensus
- hard solutions favor **Consistency**
 - 2PC, PAXOS, RAFT
 - if a node is unreachable then consensus cannot be reached
- soft solutions favor **Availability**
 - quorum, allow nodes to drop out
 - partitions can lead to more than one active set
- you **need to choose** which one is appropriate

Fighting Message Loss

- resending is the only cure
 - sender stores the message and retries
 - recipient must send acknowledgement eventually
- fighting unreliable medium is sender's obligation
 - sender's buffer must be persistent
 - resending must begin anew after a crash

At-Least-Once Delivery with Akka

```
class ALOD(other: ActorPath) extends PersistentActor with AtLeastOnceDelivery {  
    val IDs = Iterator from 1  
  
    def receiveCommand = {  
        case Command(x) =>  
            persist(ImportantResult(x)) { ir =>  
                sender() ! Done(x)  
                deliver(other, FollowUpCommand(x, IDs.next(), _))  
            }  
        case rc: ReceiptConfirmed =>  
            persist(rc) { rc => confirmDelivery(rc.deliveryId) }  
    }  
  
    def receiveRecover = {  
        case ImportantResult(x) => deliver(other, FollowUpCommand(x, IDs.next(), _))  
        case ReceiptConfirmed(_, id) => confirmDelivery(id)  
    }  
}
```

At-Least-Once Delivery with Akka

```
class ALOD(other: ActorPath) extends PersistentActor with AtLeastOnceDelivery {  
    val IDs = Iterator from 1  
  
    def receiveCommand = {  
        case Command(x) =>  
            persist(ImportantResult(x)) { ir =>  
                sender() ! Done(x)  
                deliver(other, FollowUpCommand(x, IDs.next(), _))  
            }  
        case rc: ReceiptConfirmed =>  
            persist(rc) { rc => confirmDelivery(rc.deliveryId) }  
    }  
  
    def receiveRecover = {  
        case ImportantResult(x) => deliver(other, FollowUpCommand(x, IDs.next(), _))  
        case ReceiptConfirmed(_, id) => confirmDelivery(id)  
    }  
}
```

At-Least-Once Delivery with Akka

```
class ALOD(other: ActorPath) extends PersistentActor with AtLeastOnceDelivery {
    val IDs = Iterator from 1

    def receiveCommand = {
        case Command(x) =>
            persist(ImportantResult(x)) { ir =>
                sender() ! Done(x)
                deliver(other, FollowUpCommand(x, IDs.next(), _))
            }
        case rc: ReceiptConfirmed =>
            persist(rc) { rc => confirmDelivery(rc.deliveryId) }
    }

    def receiveRecover = {
        case ImportantResult(x) => deliver(other, FollowUpCommand(x, IDs.next(), _))
        case ReceiptConfirmed(_, id) => confirmDelivery(id)
    }
}
```

At-Least-Once Delivery with Akka

```
class ALOD(other: ActorPath) extends PersistentActor with AtLeastOnceDelivery {  
    val IDs = Iterator from 1  
  
    def receiveCommand = {  
        case Command(x) =>  
            persist(ImportantResult(x)) { ir =>  
                sender() ! Done(x)  
                deliver(other, FollowUpCommand(x, IDs.next(), _))  
            }  
        case rc: ReceiptConfirmed =>  
            persist(rc) { rc => confirmDelivery(rc.deliveryId) }  
    }  
  
    def receiveRecover = {  
        case ImportantResult(x) => deliver(other, FollowUpCommand(x, IDs.next(), _))  
        case ReceiptConfirmed(_, id) => confirmDelivery(id)  
    }  
}
```

At-Least-Once Delivery with Akka

```
class ALOD(other: ActorPath) extends PersistentActor with AtLeastOnceDelivery {
    val IDs = Iterator from 1

    def receiveCommand = {
        case Command(x) =>
            persist(ImportantResult(x)) { ir =>
                sender() ! Done(x)
                deliver(other, FollowUpCommand(x, IDs.next(), _))
            }
        case rc: ReceiptConfirmed =>
            persist(rc) { rc => confirmDelivery(rc.deliveryId) }
    }

    def receiveRecover = {
        case ImportantResult(x) => deliver(other, FollowUpCommand(x, IDs.next(), _))
        case ReceiptConfirmed(_, id) => confirmDelivery(id)
    }
}
```

Mysterious Double-Bookings

- acknowledgements can be lost
- sender will faithfully duplicate the message
- therefore named **at-least-once**
- recipient is responsible for dealing with this

Exactly-Once Delivery

- processing occurs only once for each message
- all middleware does it*
 - EIP: Guaranteed Delivery just means persistence
 - JMS: configure session for AUTO_ACKNOWLEDGE
 - Google's MillWheel

Exactly-Once Delivery with Akka

```
class PersistThenAct extends PersistentActor {  
    var nextId = 1  
    def receiveCommand = {  
        case FollowUpCommand(x, id, deliveryId) =>  
            val rc = ReceiptConfirmed(id, deliveryId)  
            if (id == nextId) {  
                persist(rc) { rc =>  
                    sender() ! rc  
                    doTheAction()  
                    nextId += 1  
                }  
            } else if (id < nextId) sender() ! rc  
            // otherwise an older command was lost, so wait for the resend  
    }  
    def receiveRecover = {  
        case ReceiptConfirmed(id, _) => nextId = id + 1  
    }  
    private def doTheAction(): Unit = () // actually do the action  
}
```

Exactly-Once Delivery with Akka

```
class PersistThenAct extends PersistentActor {
    var nextId = 1
    def receiveCommand = {
        case FollowUpCommand(x, id, deliveryId) =>
            val rc = ReceiptConfirmed(id, deliveryId)
            if (id == nextId) {
                persist(rc) { rc =>
                    sender() ! rc
                    doTheAction()
                    nextId += 1
                }
            } else if (id < nextId) sender() ! rc
            // otherwise an older command was lost, so wait for the resend
    }
    def receiveRecover = {
        case ReceiptConfirmed(id, _) => nextId = id + 1
    }
    private def doTheAction(): Unit = () // actually do the action
}
```

Exactly-Once Delivery with Akka

```
class PersistThenAct extends PersistentActor {
    var nextId = 1
    def receiveCommand = {
        case FollowUpCommand(x, id, deliveryId) =>
            val rc = ReceiptConfirmed(id, deliveryId)
            if (id == nextId) {
                persist(rc) { rc =>
                    sender() ! rc
                    doTheAction()
                    nextId += 1
                }
            } else if (id < nextId) sender() ! rc
            // otherwise an older command was lost, so wait for the resend
    }
    def receiveRecover = {
        case ReceiptConfirmed(id, _) => nextId = id + 1
    }
    private def doTheAction(): Unit = () // actually do the action
}
```

Exactly-Once Delivery with Akka

```
class PersistThenAct extends PersistentActor {
    var nextId = 1
    def receiveCommand = {
        case FollowUpCommand(x, id, deliveryId) =>
            val rc = ReceiptConfirmed(id, deliveryId)
            if (id == nextId) {
                persist(rc) { rc =>
                    sender() ! rc
                    doTheAction()
                    nextId += 1
                }
            } else if (id < nextId) sender() ! rc
                // otherwise an older command was lost, so wait for the resend
    }
    def receiveRecover = {
        case ReceiptConfirmed(id, _) => nextId = id + 1
    }
    private def doTheAction(): Unit = () // actually do the action
}
```

Exactly-Once Delivery with Akka

```
class PersistThenAct extends PersistentActor {
    var nextId = 1
    def receiveCommand = {
        case FollowUpCommand(x, id, deliveryId) =>
            val rc = ReceiptConfirmed(id, deliveryId)
            if (id == nextId) {
                persist(rc) { rc =>
                    sender() ! rc
                    doTheAction()
                    nextId += 1
                }
            } else if (id < nextId) sender() ! rc
                // otherwise an older command was lost, so wait for the resend
        }
        def receiveRecover = {
            case ReceiptConfirmed(id, _) => nextId = id + 1
        }
        private def doTheAction(): Unit = () // actually do the action
    }
}
```

Exactly-Once Delivery with Akka

```
class PersistThenAct extends PersistentActor {
    var nextId = 1
    def receiveCommand = {
        case FollowUpCommand(x, id, deliveryId) =>
            val rc = ReceiptConfirmed(id, deliveryId)
            if (id == nextId) {
                doTheAction()
                persist(rc) { rc =>
                    sender() ! rc
                    nextId += 1
                }
            } else if (id < nextId) sender() ! rc
            // otherwise an older command was lost, so wait for the resend
    }
    def receiveRecover = {
        case ReceiptConfirmed(id, _) => nextId = id + 1
    }
    private def doTheAction(): Unit = () // actually do the action
}
```

The Hard Truth

CPU can stop between any two instructions, therefore exactly-once semantics cannot be Guaranteed.*

**with a capital G, i.e. with 100% absolute coverage*

Mitigating that Hard Truth

- probability of loss or duplication can be small
- many systems are fine with almost-exactly-once
- still need to decide on which side to err
 - persist-then-act → almost-exactly-but-at-most-once
 - act-then-persist → almost-exactly-but-at-least-once

We Can Do Better!

- an operation is idempotent if multiple application produces the same result as single application
- (side-)effects must also be idempotent
 - unique transaction identifiers etc.
- this achieves **effectively-exactly-once**
- can save persistence round-trip latency by optimistic execution

Effectively-Exactly-Once with Akka

```
class Idempotent(other: ActorRef) extends PersistentActor {
  var nextId = 1
  def receiveCommand = {
    case FollowUpCommand(x, id, deliveryId) =>
      val rc = ReceiptConfirmed(id, deliveryId)
      if (id == nextId) {
        // optimistically execute
        if (isValid(x)) {
          updateState(x)
          other ! Command(x) // assuming idempotent receiver
        } else sender() ! Invalid(x)
        nextId += 1
        // then take care of reliable delivery handling
        persistAsync(rc) { rc =>
          sender() ! rc
        }
      } else if (id < nextId) sender() ! rc
    }
    ...
}
```

The Future is Bright Again!



source: <http://freewallpapersbackgrounds.com/>

Consistency à la Carte

- not all parts of a system have the same needs
- expend the effort only where necessary
 - idempotency does not come for free
 - confirmation handling cannot always be abstracted away
- incur the runtime overhead only where needed
 - persistence round-trips take time
 - consensus protocols take more time
- include required semantics in system design

**Remember that the trade-off is always between
consistency and availability.**

And latency.



©Typesafe 2014 – All Rights Reserved